

## ▼ Task

Implement Linear and Logistic Regression on real-world datasets. Implementing Linear Regression

### ▼ Upload file

Subtask:

Create a cell for uploading the csv files.

**Reasoning:** Create a code cell to handle the file upload process using google.colab.files.upload and print a message to the user.

```
# Linear Regression on California Housing Dataset

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score

# -----
# 1. Load dataset
# -----
df = pd.read_csv("housing.csv")

print(df.head())
print(df.info())

# -----
# 2. Handle missing values
# -----
df['total_bedrooms'] = df['total_bedrooms'].fillna(df['total_bedrooms'].median())

# -----
# 3. Features & target
# -----
X = df.drop("median_house_value", axis=1)
y = df["median_house_value"]

# -----
# 4. Identify column types
# -----
numeric_features = X.select_dtypes(include=[np.number]).columns
categorical_features = ['ocean_proximity']

# -----
# 5. Preprocessing pipeline
# -----
numeric_transformer = Pipeline(steps=[
    ("scaler", StandardScaler())
])

categorical_transformer = Pipeline(steps=[
    ("onehot", OneHotEncoder(handle_unknown="ignore"))
])

preprocessor = ColumnTransformer(
    transformers=[
        ("num", numeric_transformer, numeric_features),
        ("cat", categorical_transformer, categorical_features)
    ]
)

# -----
# 6. Train-test split
# -----
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
```

```
# -----
# 7. Linear Regression model
# -----
model = Pipeline(steps=[
    ("preprocessor", preprocessor),
    ("regressor", LinearRegression())
])

# -----
# 8. Train model
# -----
model.fit(X_train, y_train)

# -----
# 9. Predictions
# -----
y_pred = model.predict(X_test)

# -----
# 10. Evaluation
# -----
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

print(f"RMSE: {rmse:.2f}")
print(f"R2 Score: {r2:.4f}")

# -----
# 11. Visualization
# -----
plt.figure(figsize=(6,6))
plt.scatter(y_test, y_pred, alpha=0.3)
plt.xlabel("Actual House Value")
plt.ylabel("Predicted House Value")
plt.title("Linear Regression: Actual vs Predicted")
plt.plot([y_test.min(), y_test.max()],
         [y_test.min(), y_test.max()],
         color="red")
plt.show()
```

```

longitude  latitude  housing_median_age  total_rooms  total_bedrooms \
0   -122.23    37.88           41.0      880.0        129.0
1   -122.22    37.86           21.0      7099.0       1106.0
2   -122.24    37.85           52.0      1467.0        190.0
3   -122.25    37.85           52.0      1274.0       235.0
4   -122.25    37.85           52.0      1627.0       280.0

population  households  median_income  median_house_value  ocean_proximity
0      322.0        126.0      8.3252      452600.0      NEAR BAY
1     2401.0       1138.0      8.3014      358500.0      NEAR BAY
2      496.0        177.0      7.2574      352100.0      NEAR BAY
3      558.0        219.0      5.6431      341300.0      NEAR BAY
4      565.0        259.0      3.8462      342200.0      NEAR BAY

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 10 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   longitude         20640 non-null   float64
 1   latitude          20640 non-null   float64
 2   housing_median_age 20640 non-null   float64
 3   total_rooms       20640 non-null   float64
 4   total_bedrooms    20640 non-null   float64
 5   population        20640 non-null   int64  
 6   households        20640 non-null   int64  
 7   median_income     20640 non-null   float64
 8   median_house_value 20640 non-null   float64
 9   ocean_proximity   20640 non-null   object  
dtypes: float64(9), object(1)
memory usage: 1.6+ MB
None
RMSE: 70060.52
R2 Score: 0.6254

```

